

Sanmi Koyejo

Assistant Professor
Department of Computer Science,
Stanford University.

Mail: Room 342, 353 Jane Stanford Way,
Stanford, CA 94305
Email: sanmi@stanford.edu
Web: <http://cs.stanford.edu/~sanmi>

Academic Positions

2022 -	Assistant Professor	Stanford University
2022 -	Adjunct Associate Professor	University of Illinois at Urbana-Champaign
2021 - 2022	Associate Professor	University of Illinois at Urbana-Champaign
2016 - 2021	Assistant Professor	University of Illinois at Urbana-Champaign
2017 - 2017	Visiting Scholar	UC Berkeley (J.B. Poline & M. D'Esposito)
2014 - 2016	Engineering Research Associate	Stanford University (PI: R. Poldrack)
2013 - 2014	Research Associate	UT Austin (PI: R. Poldrack & P. Ravikumar)

Education

2013 - 2013	Postdoctoral Fellow	UT Austin (PI: R. Poldrack & P. Ravikumar)
2013	Ph.D., Electrical Engineering	UT Austin (Advisor: J. Ghosh)
2008	M.S., Electrical Engineering	UT Austin
2005	B.S. (Hons), Electrical Engineering & Minor in Statistics	New Jersey Institute of Technology (NJIT)

Recent Awards

Presidential Early Career Award for Scientists and Engineers (PECASE) Award (2025), NeurIPS Outstanding Paper Award (2023) x2, Terman Faculty Fellowship (2022); Skip Ellis Early Career Award (2021); Alfred P. Sloan Research Fellow in Computer Science (2021); NSF CAREER Award (2021); Early Career Spotlight (IJCAI, 2019); Kavli Fellow (2018); UAI Outstanding Paper Award (2013)

Recent Service

Black in AI President (2021 -); Association for Health, Learning and Inference; Board of Directors (2021 -); Neural Information Processing Systems Foundation; Board of Directors (2022 -); Co-General Chair, Neural Information Processing Systems Conference (NeurIPS, 2022); Action Editor, Journal of Machine Learning Research (2021 -); Board of Reviewing Editors, eLife (2020-2023); Workshop Chair, Neural Information Processing Systems Conference (NeurIPS, 2021); Diversity and Inclusion Co-chair, Artificial Intelligence and Statistics (AISTATS 2021); Workshop Co-chair, International Conference on Learning Representations (ICLR, 2021)

Selected Talks & Tutorials

Measurement and Mismeasurement of AI (FAccT 2024); Towards Trustworthy LLMs (WSDM 2024); Algorithmic Fairness: Why It's Hard and Why It's Interesting (CVPR 2022 Tutorial, ISBI 2023 Keynote); Representation Learning and Fairness (NeurIPS 2019 Tutorial)

Selected Publications

(Full List: <https://scholar.google.com/citations?user=Eaa0eJwAAAAJ>)

1. Sang T Truong, Duc Q Nguyen, Toan Nguyen, Dong D Le, Nhi N Truong, Tho Quan, and Sanmi Koyejo. Crossing linguistic horizons: Finetuning and comprehensive evaluation of vietnamese large language models. *North American Chapter of the Association for Computational Linguistics - Human Language Technologies (NAACL-HLT)*, 2024
2. Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Zinan Lin, Yu Cheng, Sanmi Koyejo, Dawn Song, and Bo Li. DecodingTrust: A comprehensive assessment of trustworthiness in GPT models. In *Thirty-seventh Conference on Neural Information Processing Systems; Datasets and Benchmarks Track*, 2023
3. Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. Are emergent abilities of large language models a mirage? *Advances in Neural Information Processing Systems*, 36, 2023
4. Karen Drukker, Weijie Chen, Judy Gichoya, Nicholas Gruszauskas, Jayashree Kalpathy-Cramer, Sanmi Koyejo, Kyle Myers, Rui C. Sá, Berkman Sahiner, Heather Whitney, Zi Zhang, and Maryellen Giger. Toward fairness in artificial intelligence for medical image analysis: identification and mitigation of potential biases in the roadmap from data collection to model deployment. *Journal of Medical Imaging*, 10(6):061104–061104, 2023
5. Zachary Robertson, Hantao Zhang, and Sanmi Koyejo. Cooperative inverse decision theory for uncertain preferences. In *International Conference on Artificial Intelligence and Statistics*, pages 5854–5868. PMLR, 2023
6. Jessica Schrouff, Natalie Harris, Oluwasanmi Koyejo, Ibrahim Alabdulmohsin, Eva Schnider, Krista Opsahl-Ong, Alex Brown, Subhrajit Roy, Diana Mincu, Christina Chen, Awa Dieng, Yuan Liu, Vivek Natarajan, Alan Karthikesalingam, Katherine Heller, Silvia Chiappa, and Alexander D’Amour. Diagnosing failures of fairness transfer across distribution shift in real-world medical settings. *Advances in Neural Information Processing Systems*, 35:19304–19318, 2022
7. Yong Xie, Dakuo Wang, Pin-Yu Chen, Jinjun Xiong, Sijia Liu, and Oluwasanmi Koyejo. A word is worth a thousand dollars: Adversarial attack on tweets fools stock prediction. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 587–599. Association for Computational Linguistics, July 2022
8. Edited by: Peter Kairouz and H. Brendan McMahan. Advances and open problems in federated learning. *Foundations and Trends® in Machine Learning*, 14(1):–, 2021
9. James M Shine, Oluwasanmi Koyejo, and Russell A Poldrack. Temporal metastates are associated with differential patterns of time-resolved connectivity, network topology, and attention. *Proceedings of the National Academy of Sciences*, 2016
10. Xiaojun Xu, Jacky Y Zhang, Evelyn Ma, Hyun Ho Son, Sanmi Koyejo, and Bo Li. Adversarially robust models may not transfer better: Sufficient conditions for domain transferability from the view of regularization. In *International Conference on Machine Learning*, pages 24770–24802. PMLR, 2022

References

- [1] Edited by: Peter Kairouz and H. Brendan McMahan. Advances and open problems in federated learning. *Foundations and Trends® in Machine Learning*, 14(1):–, 2021.
- [2] Karen Drukker, Weijie Chen, Judy Gichoya, Nicholas Gruszauskas, Jayashree Kalpathy-Cramer, Sanmi Koyejo, Kyle Myers, Rui C. Sá, Berkman Sahiner, Heather Whitney, Zi Zhang, and Maryellen Giger. Toward fairness in artificial intelligence for medical image analysis: identification and mitigation of potential biases in the roadmap from data collection to model deployment. *Journal of Medical Imaging*, 10(6):061104–061104, 2023.
- [3] Zachary Robertson, Hantao Zhang, and Sanmi Koyejo. Cooperative inverse decision theory for uncertain preferences. In *International Conference on Artificial Intelligence and Statistics*, pages 5854–5868. PMLR, 2023.
- [4] Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. Are emergent abilities of large language models a mirage? *Advances in Neural Information Processing Systems*, 36, 2023.
- [5] Jessica Schrouff, Natalie Harris, Oluwasanmi Koyejo, Ibrahim Alabdulmohsin, Eva Schnider, Krista Opsahl-Ong, Alex Brown, Subhrajit Roy, Diana Mincu, Christina Chen, Awa Dieng, Yuan Liu, Vivek Natarajan, Alan Karthikesalingam, Katherine Heller, Silvia Chiappa, and Alexander D’Amour. Diagnosing failures of fairness transfer across distribution shift in real-world medical settings. *Advances in Neural Information Processing Systems*, 35:19304–19318, 2022.
- [6] James M Shine, Oluwasanmi Koyejo, and Russell A Poldrack. Temporal metastates are associated with differential patterns of time-resolved connectivity, network topology, and attention. *Proceedings of the National Academy of Sciences*, 2016.
- [7] Sang T Truong, Duc Q Nguyen, Toan Nguyen, Dong D Le, Nhi N Truong, Tho Quan, and Sanmi Koyejo. Crossing linguistic horizons: Finetuning and comprehensive evaluation of vietnamese large language models. *North American Chapter of the Association for Computational Linguistics - Human Language Technologies (NAACL-HLT)*, 2024.
- [8] Boxin Wang, Weixin Chen, Hengzhi Pei, Chulin Xie, Mintong Kang, Chenhui Zhang, Chejian Xu, Zidi Xiong, Ritik Dutta, Rylan Schaeffer, Sang T. Truong, Simran Arora, Mantas Mazeika, Dan Hendrycks, Zinan Lin, Yu Cheng, Sanmi Koyejo, Dawn Song, and Bo Li. DecodingTrust: A comprehensive assessment of trustworthiness in GPT models. In *Thirty-seventh Conference on Neural Information Processing Systems; Datasets and Benchmarks Track*, 2023.
- [9] Yong Xie, Dakuo Wang, Pin-Yu Chen, Jinjun Xiong, Sijia Liu, and Oluwasanmi Koyejo. A word is worth a thousand dollars: Adversarial attack on tweets fools stock prediction. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 587–599. Association for Computational Linguistics, July 2022.
- [10] Xiaojun Xu, Jacky Y Zhang, Evelyn Ma, Hyun Ho Son, Sanmi Koyejo, and Bo Li. Adversarially robust models may not transfer better: Sufficient conditions for domain transferability from the view of regularization. In *International Conference on Machine Learning*, pages 24770–24802. PMLR, 2022.