# Brando Miranda, Alycia Lee

Dept. of Computer Science
Advised by Sanmi Koyejo, Stanford Computer Science (CS)
In Collaboration with Patrick Yu


Brando's website

# The Diversity Coefficient: A Data Quality Metric that shows LLMs are Pretrained on Formally Diverse Data


*Keywords: large language models (LLMs), data quality, metrics, diversity*


## Summary


- We develop a data quality metric to measure the formal diversity of the pretraining data of Large Language Models (LLMS).
- Diversity coefficient of LLM pretraining datasets are high compared to previous work on vision datasets.
- We test that the diversity coefficient correlates with the ground truth diversity (when known).
- The diversity coefficient passes important sanity checks:
  - The diversity coefficient correlates with latent concepts and vocab size in the synthetic GINC (language) datasets.
  - The diversity coefficient estimates low diversity when comparing tasks of the same data set, but high diversity in cross task comparisons is across different data sets.


## More about Brando Miranda

- Current EDGE Scholar at Stanford University.
- Research interests lie in meta-learning, foundation models for theorem proving, and human & brain inspired AI.
- M.Eng. in Electrical Engineering and Computer Science at MIT.

Stanford | Data Science


NSF III 2046795, IIS 1909577, CCF 1934986, NIH 1R01MH116226-01A, NIFA award 2020-67021-32799, the Alfred P. Sloan Foundation, and Google Inc.


## Methods

**Task2Vec-based diversity coefficient** (Miranda, Yu, et al 2022) approximately measures the intrinsic variability of tasks in a few-shot learning benchmark.

Ground Truth Diversity Coefficient: expected distance between pairs of tasks T1, T2:

$$\hat{div}(B) = \mathbb{E}_{\tau_1,\tau_2\sim\hat{p}(\tau|B):\tau_1\neq\tau_2}\left[d(p(x_1,y_1 \mid \tau_1), p(x_2,y_2 \mid \tau_2))\right]$$

Definition: expected distance between pairs of tasks T1, T2 as Task2Vec embeddings:

$$\hat{div}(B) = \mathbb{E}_{\tau_1,\tau_2\sim\hat{p}(\tau|B):\tau_1\neq\tau_2}\mathbb{E}_{D_1\sim\hat{p}(x_1,y_1|\tau_1),D_2\sim\hat{p}(x_2,y_2|\tau_2)}\left[d(\mathrm{diag}(\hat{F}_{D_1,f_w}), \mathrm{diag}(\hat{F}_{D_2,f_w})\right]$$

## Experiments & Results

**Task2Vec diversity coefficient correlates with ground truth diversity for synthetic data**

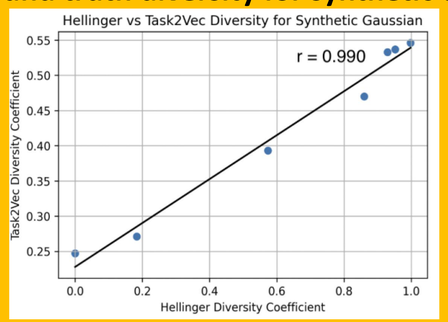


**Diversity of c4, Wikitext-103, and The Pile are twice as high vs. vision benchmarks**

| Dataset | Probe Network | Diversity Coeff |
|---|---|---|
| MiniImagenet | Resnet18 | 0.117 ±2.098e-5 |
| Cifar-fs | Resnet18 | 0.100 ±2.18e-5 |
| c4 | GPT-2 | **0.2374 ±2.785e-5** |
| WikiText-103 | GPT-2 | **0.2140 ±7.93e-5** |
| The Pile | GPT-2 | **0.2463 ±3.034-5** |

**Batch size correlates with diversity**

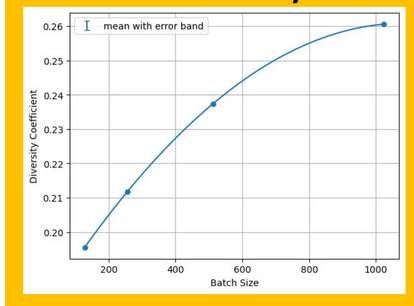


**Random probe underestimates diversity, non fine-tuned overestimates diversity**

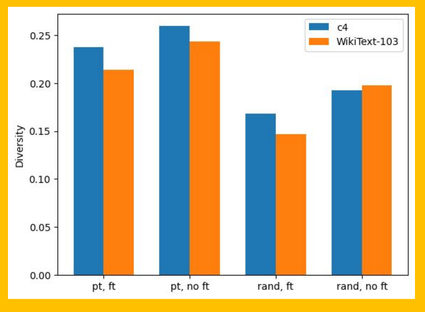


**Pairwise combinations of The Pile datasets have higher diversity vs. individual datasets**

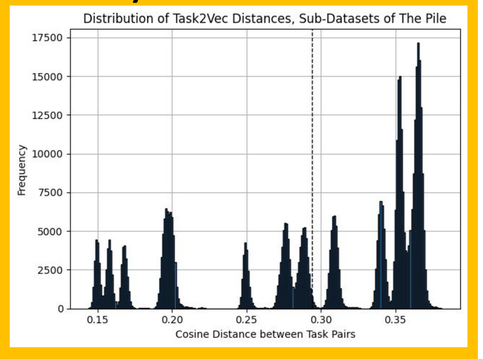


**Diversity correlates with # latent concepts (left) and vocab size (right) in GINC**