

brando9@stanford.edu
Stanford, CA, 94305.

Brando Miranda

Gates Bldg., 353 Jane Stanford Way

Google Scholar; Website; Stanford Profile; MIT Website

RESEARCH INTERESTS

AI for formal mathematics and Lean 4 theorem proving; autoformalization, end-to-end formal verification, and contamination-resistant evaluation of LLM mathematical reasoning; data-centric machine learning for foundation models.

EDUCATION

Ph.D. in Computer Science

Stanford University

Advisor: Prof. Sanmi Koyejo, Stanford Trustworthy AI Research Group (STAIR).

2022-2026 (*expected*)

GPA: 4.045/4.0

Master of Engineering in Electrical Engineering and Computer Science

Massachusetts Institute of Technology

Advisor: Prof. Tomaso Poggio, Center for Brains, Minds and Machines.

Thesis Title: *Function Approximation with Deep Neural and Gaussian Networks*.

2014-2016

GPA: 4.8/5.0

Bachelor of Science, Computer Science and Engineering

Massachusetts Institute of Technology

2010-2014

Minors: Mathematics, Music

AWARDS & HONORS

- **ICML 2026 Silver Reviewer** (May 2026) — top reviewer recognition, signed by the ICML 2026 Program Chairs.
- **NeurIPS Outstanding Main Track Paper Award** (December 2023) — top 0.4% of NeurIPS submissions; only 2 main-track papers selected. For “Are Emergent Abilities of Large Language Models a Mirage?” (Schaeffer, Miranda, Koyejo).
- ICML Workshop on Trustworthy Multi-modal Foundation Models and AI Agents (TiFA) — Outstanding Paper Award (July 2024) — for “Why Has Predicting Downstream Capabilities of Frontier AI Models with Scale Remained Elusive?”
- **EDGE Scholar**, Stanford University (September 2022) — Stanford fellowship supporting first-generation / low-income PhD scholars.
- **Stanford School of Engineering Fellowship** (September 2022) — multi-year departmental fellowship for incoming engineering PhDs.
- **Honorable Mention**, Ford Foundation Predoctoral Fellowship (2020, 2021) — national fellowship recognizing PhD applicants with potential to diversify the U.S. professoriate.
- **Best Research Project Award**, UIUC graduate course CS 598 “Learning to Learn” (December 2020).
- **Most Cited Paper Certificate**, International Journal of Automation & Computing (IJAC, December 2019) — for “Why and when can deep- but not shallow-networks avoid the curse of dimensionality: a review”.
- **Computer Science Excellence Saburo Muroga Endowed Fellow**, UIUC (2019-2020) — top departmental fellowship for incoming CS PhDs.
- **Sloan Scholar**, Alfred P. Sloan Foundation Minority Ph.D. (MPHD) Program (2018-2019) — national fellowship for underrepresented STEM PhD students.
- **Grainger Engineering SURGE Fellowship**, UIUC (2018-2019) — multi-year college-level fellowship for diverse engineering PhDs.

PUBLICATIONS

Refereed Publications

(1) *VeriBench-FTP: A Formal Theorem Proving Benchmark in Lean 4 for Code Verification* (2025)

S. Barkallah, S. Daruru, Brando Miranda, Leni Aniva, Anima Nie, Sanmi Koyejo

5th NeurIPS Workshop on Mathematical Reasoning and AI. 2025

(2) *VeriBench: End-to-End Formal Verification Benchmark for AI Code Generation in Lean 4* (2025)

Brando Miranda, Zhangir Zhou, Anima Nie, Elio Obbad, Leni Aniva, Kai Fronsdal, William Kirk, Dogus Soylu, et al.

2nd Workshop on AI for Math at International Conference on Machine Learning (ICML). 2025

[OpenReview]

(3) *Pantograph: A machine-to-machine interaction interface for advanced theorem proving, high level reasoning, and data extrac* (2025)

Leni Aniva, Chuyue Sun, Brando Miranda, Clark Barrett, Sanmi Koyejo

International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS). 2025 —
International Conference on Tools and Algorithms for the Construction and Analysis of Systems
[Springer Link]

(4) *Putnam-AXIOM: A Functional & Static Benchmark for Measuring Higher Level Mathematical Reasoning in LLMs* (2025)

Aryan Gulati, Brando Miranda, Eric Chen, Emily Xia, Kai Fronsdal, Bruno de Moraes Dumont, Sanmi Koyejo
International Conference on Machine Learning (ICML). 2025 & NeurIPS Workshop on Mathematics and AI (MATH-AI).
2024
[arXiv]

(5) *Why Has Predicting Downstream Capabilities of Frontier AI Models with Scale Remained Elusive?* (2025)

Rylan Schaeffer, Hailey Schoelkopf, Brando Miranda, Gabriel Mukobi, Varun Madan, Adam Ibrahim, Herbie Bradley,
Stella Biderman, Sanmi Koyejo
International Conference on Machine Learning (ICML). 2025 & ICML Workshop on Trustworthy Multi-modal Foundation
Models and AI Agents (TiFA) — Outstanding Paper Award. 2024
[arXiv]

(6) *CoDaPO: Confidence and Difficulty-Adaptive Policy Optimization for Post-Training Language Models* (2025)

Zhangir Zhou, Xingrun Lu, Cheng Cao, Brando Miranda, Tong Liu, Bo Han, Sanmi Koyejo
2nd Workshop on AI for Math at International Conference on Machine Learning (ICML). 2025
[OpenReview]

(7) *Causally Quantifying the Effect of Test Set Contamination on Generative Benchmarks* (2025)

Rylan Schaeffer, Brando Miranda, Jiaqi Kazdan, Kaivu Liu, Amir M. Ahmed, Niloofar Miresghallah, et al.
NeurIPS Workshop on Evaluating the Evolving LLM Lifecycle. 2025

(8) *Position: Machine Learning Conferences Should Establish a ‘Refutations and Critiques’ Track* (2025)

Rylan Schaeffer, Jiaqi Kazdan, Yarin Denirov-Blanch, Brando Miranda, Max Gerstgrasser, et al.
Neural Information Processing Systems (NeurIPS), Main Track. 2025
[arXiv]

(9) *Failures to Find Transferable Image Jailbreaks Between Vision-Language Models* (2024)

Rylan Schaeffer, Dan Valentine, Luke Bailey, James Chua, et al., Brando Miranda, et al., Sanmi Koyejo, Ethan Perez
International Conference on Learning Representations (ICLR). 2024 & NeurIPS Workshop on Red Teaming Generative AI,
Contributed Talk. 2024
[OpenReview]

(10) *Does Maximizing Neural Regression Scores Teach Us About The Brain?* (2024)

Rylan Schaeffer, Mikail Khona, Sankarshan Chandra, Michael Ostrow, Brando Miranda, Sanmi Koyejo
NeurIPS Workshop on Unifying Representations in Neural Models (UniReps), 2nd Edition. 2024
[OpenReview]

(11) *Quantifying the Importance of Data Alignment in Downstream Model Performance* (2024)

Keshav Chawla, Aditya Sahai, Marco DePavia, Sathvik Sundar, Brando Miranda
ICLR Workshop on Data-Centric Machine Learning Research (DMLR). 2024 (DMLR)
[arXiv]

(12) *An Evaluation Benchmark for Autoformalization in Lean4* (2024)

Aryan Gulati, Devanshu Ladsaria, Shubham Mishra, Jasdeep Sidhu, Brando Miranda
ICLR Tiny Papers Track, Second Edition. 2024
[arXiv]

(13) *Are Emergent Abilities of Large Language Models a Mirage?* (2023)

Rylan Schaeffer, Brando Miranda, Sanmi Koyejo
NeurIPS Outstanding Main Track Paper Award & Oral 2023
**Featured in: New York Times, Quanta Magazine, Forbes, Stanford HAI; Cited in 2024 Economic Report
of the President**
[arXiv] [OpenReview]

(14) *Beyond Scale: the Diversity Coefficient as a Data Quality Metric Demonstrates LLMs are Pre-trained on Formally Diverse*
(2023)

*Alycia Lee**, *Brando Miranda**, *Patrick Yu*, *Oluwasanmi Koyejo* (*equal contribution)

ICML Workshop on Data-Centric Machine Learning & ICML Workshop on Deployable Generative AI. 2023

[arXiv]

(15) *Why and when can deep-but not shallow-networks avoid the curse of dimensionality: a review* (2017)

Tomaso Poggio, Hrushikesh Mhaskar, Lorenzo Rosasco, Brando Miranda, Qianli Liao

International Journal of Automation and Computing 2017, Most Cited Paper Certificate

[Journal]

Preprints and Technical Reports

(1) *AI Coding Benchmarks Need Proofs, Not Just Tests* (2026)

D. Amrollahi, M. Karimi, Brando Miranda, Leni Aniva, Chuyue Sun, Clark Barrett, Sanmi Koyejo

Preprint 2026 — **Lean 4** formal-proof-based evaluation of AI-generated code

[PDF]

(2) *Lean-ing on Quality: How High-Quality Data Beats Diverse Multilingual Data in Autoformalization* (2025)

W. Chan, M. Souliman, J. Nordhagen, Brando Miranda, Elyas Obbad, Sanmi Koyejo

Preprint, arXiv:2502.15795

[arXiv]

(3) *ZIP-FIT: Embedding-Free Data Selection via Compression-Based Alignment* (2024)

Elyas Obbad, Iddah Mlauzi, Brando Miranda, Rylan Schaeffer, Kamal Obbad, Suhana Bedi, Sanmi Koyejo

Preprint, arXiv:2410.18194

[arXiv] [Code]

PROFESSIONAL EXPERIENCE

Stanford University - Stanford, CA <i>Ph.D. Student in Computer Science. Advisor: Professor Sanmi Koyejo</i>	<i>September 2022 - 2026 (expected)</i>
Amazon Web Services (AWS) - Cupertino, CA <i>Applied Scientist Intern</i>	<i>June 2024 - September 2024</i>
Morph Labs - Remote <i>Machine Learning Research Scientist Consultant</i>	<i>October 2023 - December 2023</i>
Wise Agents - Stanford Spin-out <i>AI Research Consultant</i>	<i>2023</i>
IBM Research - Yorktown Heights, NY <i>Graduate Research Intern</i>	<i>May 2022 - August 2022</i>
University of Illinois Urbana-Champaign - Urbana-Champaign, IL <i>Ph.D. Student in Computer Science. Advisor: Professor Sanmi Koyejo</i>	<i>September 2018 - May 2022</i>
IBM Research - Yorktown Heights, NY <i>Graduate Research Intern</i>	<i>May 2021 - August 2021</i>
MIT CBMM (Center for Brain Minds & Machines) - Cambridge, MA <i>Research Assistant. Advisor: Professor Tomaso Poggio</i>	<i>June 2015 - September 2018</i>

MEDIA COVERAGE

- **White House Economic Report of the President (2024)**: Our work on emergent abilities was cited in the 2024 Economic Report of the President – the White House’s annual flagship document on national economic policy
- **The New York Times (2023)**: "Silicon Valley Confronts the Idea That the 'Singularity' Is Here"
- **Quanta Magazine (2024)**: "How Quickly Do Large Language Models Learn Unexpected Skills?"
- **Forbes (2023)**: "AI 'Emergent Abilities' Are A Mirage, Says AI Researcher"
- **Andrew Ng (2024)**: Endorsed our paper as evidence that AGI development will be smooth and predictable
- **Hacker News / Y Combinator (2025)**: Putnam-AXIOM benchmark posted on Hacker News — our Putnam-AXIOM benchmark for higher-level LLM mathematical reasoning drew front-page attention on Y Combinator’s tech-news forum
- **Additional coverage**: Stanford HAI, Y Combinator News, Vice, Medium, HackerNews, NeurIPS blog, Reddit

INVITED & CONTRIBUTED TALKS

- **Applications of Trustworthy Machine Reasoning with AI Coding Agents** – AAAI 2026 Tutorial "Trustworthy Machine Reasoning with Foundation Models" (Part IV, 50 min), Singapore, January 2026
- **Emergent Abilities in Large Language Models: Mirage or an Elusive Predictive Frontier?** – Hong Kong Baptist University, Department of Computer Science Seminar, Hong Kong, January 2026
- **Are Emergent Abilities of Large Language Models a Mirage?** – Stanford IEEE Invited Talk, Stanford, CA, 2023
- **Failures to Find Transferable Image Jailbreaks Between Vision-Language Models** – NeurIPS Red Teaming GenAI Workshop, Contributed Talk, December 2024
- **The Curse of Low Task Diversity: On the Failure of Transfer Learning to Outperform MAML** – NeurIPS Meta-Learning Workshop, Contributed Talk, New Orleans, LA, December 2022

TEACHING EXPERIENCE

University of Illinois Urbana-Champaign <i>Graduate Teaching Assistant</i> • CS 446 Machine Learning • Designed problem sets and exams; held weekly office hours	<i>August 2020 - December 2020</i>
Massachusetts Institute of Technology <i>Graduate Teaching Assistant</i> • Statistical Learning Theory & Applications (9.520/6.860) with Prof. Lorenzo Rosasco & Tomaso Poggio • Introduction to Algorithms (6.006) with Prof. Nancy Lynch, Bruce Tidor and Aleksander Madry • Design & Analysis of Algorithms (6.046) with Prof. Shafi Goldwasser, Dana Moshkovitz, Nir Shavit • Introduction to Machine Learning (6.036) with Prof. Tommi Jaakkola, Suvrit Sra & Regina Barzilay	<i>2014-2016</i>

LEADERSHIP & SERVICE

Research Mentorship

2016 - Present

- Mentored undergraduate and graduate students at Stanford, UIUC, and MIT CBMM
- Led research teams on data quality metrics, meta-learning, and formal reasoning

Academic Service

2018 - Present

- Reviewer for ICML 2026 (**Silver Reviewer** – top reviewer recognition), ICLR 2020, JMLR 2018
- Graduate advisor for Latinos in Computer Science (LCS) at UIUC
- Founding member of Stanford AI for Lean, a research community advancing AI for Lean theorem proving and formalizing mathematics
- Founded "Stanford Bachata Sensual & Brazilian Zouk" and "UIUC Bachata Sensual & Zouk" official student organizations

Outreach

2017 - Present

- Distributed Research Experiences for Undergraduates (DREU) UIUC 2019
- Undergraduate Research Opportunity Program (UROP) MIT 2017-2018
- Engineering of Intelligence Team (EIT) CBMM MIT 2017-2018

SKILLS & TOOLS

Programming: Python, PyTorch, TensorFlow, JAX, Lean 4, OCaml, Coq

ML/AI Tools: Hugging Face, Weights & Biases, DSPy, git, UNIX

Languages: English, Spanish (native fluency)